



Querying LLMs with Retrieved Data

1. Overview of AI Workflows	2. Why Postgres as a vector store	3. Storing and managing vectors	4. Querying the vector store
<ul style="list-style-type: none"> Look at high-level architecture - LLMs, vector stores and JSON Look at key vocabulary and concepts (embeddings, vectors, hybrid queries, etc.) 	<ul style="list-style-type: none"> What is a vector store? Key concepts and use cases. Why Postgres and how does it compare with other market tools Setting up Postgres with vector capabilities (pgvector) Lab: Install and configure Postgres using Docker 	<ul style="list-style-type: none"> Generating embeddings: Overview of tools and workflows Storing and organizing embeddings in Postgres Strategies for handling large datasets including chunking Dense and sparse vectors Lab: Generate embeddings for a dataset and store them 	<ul style="list-style-type: none"> Techniques for similarity search: k-NN, cosine similarity Using indexes to optimize vector queries Reranking results Lab: Query stored vectors to retrieve similar items (document/image search)
5. Querying LLMs with retrieved data	6. NoSQL with JSON in Postgres	7. Integrating Vector, Relational and JSON Data	8. Putting it all together
<ul style="list-style-type: none"> Recap on querying LLMs vis APIs Best practices for combining vector retrieval with LLM prompts Prompt configuration parameters (temperature, top-k, etc) Lab: Build a pipeline where vector store results enhance LLM responses (context-aware Q&A, etc) 	<ul style="list-style-type: none"> Overview of JSON/JSONB support in Postgres Querying JSONB data with SQL Indexing JSONB data for performance Lab: Design a schema mixing vector, relational and JSONB data for a sample project 	<ul style="list-style-type: none"> Building hybrid queries to power advanced workflows Case study: Combining embeddings, metadata (relational) and configurations (JSON) Lab: Implement a hybrid query to support a sample AI use case 	<ul style="list-style-type: none"> Full stack pipeline demo: Retrieve data, query the LLM and return results Debugging and optimising the workflow Spotlight on LLM frameworks Lab: Build a working application combining all elements



Plan for the session

- Understand the Fundamentals of LLM Querying
- Apply Best Practices for Combining Vector Retrieval with LLM Prompts
- Optimize LLM Queries with Configuration Parameters
- Lab: **Build a Pipeline for Context-Aware LLM Responses**

How do we interact with LLMs?

4

- Large Language Models (LLMs) generate text based on input prompts
- They rely on probabilistic word prediction using pre-trained embeddings
- Two common ways to interact with LLMs:
 - Direct Input: User provides a prompt, LLM generates an output
 - Retrieval-Augmented Generation (RAG): LLM is provided with external data to refine responses

✦ **Key Takeaway:** LLMs do not “remember” external data unless explicitly provided in context.

Using APIs to Query LLMs

5

- Most production use cases rely on **LLM APIs** rather than self-hosted models.
- APIs provide:
 - **Standardized endpoints** (e.g., POST /v1/completions)
 - **Customization parameters** (temperature, max_tokens, top-k)
 - **Integration with retrieval pipelines**
- Example providers:
 - OpenAI (gpt-4, gpt-3.5-turbo)
 - Cohere, Anthropic (Claude), Google (Gemini)
 - Open-source models (via Hugging Face Inference API)

Demo



Common Challenges in LLM Querying

7

- **Context Length Limitations:** Most LLMs can only process a fixed number of tokens (~4K-128K).
 - **Hallucinations:** LLMs may generate plausible but incorrect information.
 - **Deterministic vs. Stochastic Behavior:**
 - Lower temperature = more predictable responses
 - Higher temperature = more creative but less reliable responses
 - **Lack of Domain Knowledge:** LLMs may not have up-to-date or domain-specific knowledge unless retrieved externally.
- 📌 **Key Takeaway:** To improve reliability, external knowledge should be retrieved and provided as context.

Enhancing LLM Queries with Retrieved Data (RAG)

8

- **What is RAG?**
 - Retrieves external data before querying an LLM.
 - Provides additional knowledge in the prompt.
 - **Steps in RAG-Based Querying:**
 - **User Query** → **Vector Search**: Find relevant stored data.
 - **Retrieve Results** → **Context Injection**: Add retrieved text to the prompt.
 - **Query LLM with Enhanced Context** → Improved response accuracy.
 - **Example Use Case:**
 - Searching an internal knowledge base before asking an LLM.
 - AI-powered search assistants retrieving FAQs before generating answers.
- 📌 **Key Takeaway:** Combining vector retrieval with LLMs ensures **more accurate, relevant, and reliable** responses.

Demo



Optimizing Vector-Augmented LLM Queries

10

✅ **Retrieve Before You Generate**

- Always fetch relevant data first, then pass it into the LLM.

✅ **Limit the Number of Retrieved Items**

- Avoid overloading the LLM with too much context (3-5 items max).

✅ **Format Context in a Structured Way**

- Use **bullet points, JSON, or key-value pairs** in the LLM prompt.

✅ **Use Clear Prompt Engineering**

- Tell the LLM exactly how to use the retrieved data (e.g., *"Summarize the books and categorize them by experience level"*).

📌 **Key Takeaway:**

The quality of LLM responses depends on how well the vector-retrieved data is structured and provided.

Problem	Cause	Fix
LLM Hallucinates Facts	LLM generates responses without sufficient grounding in real data.	Ensure retrieved data is structured before sending it to the LLM. Clearly separate facts from instructions.
Retrieved Results Are Not Relevant	Poor-quality embeddings or incorrect similarity metric used.	Use a better embedding model (bge-m3), filter by metadata , and fine-tune retrieval parameters (top-k, similarity threshold) .
LLM Response Is Too Generic	The LLM is responding broadly instead of using retrieved context.	Provide explicit instructions in the prompt , e.g., <i>"Base your answer only on the books provided."</i>
Query is Slow	Query is Slow	Query is Slow

Confidential & Proprietary – Not for Distribution

Tuning LLM Responses with Parameters

12

Key Points:

- LLM outputs can be controlled by adjusting key parameters.
- The same query can produce **deterministic, creative, or factually rich responses** depending on settings.
- Understanding parameter effects allows **fine-tuning for optimal results**.

Example Use Case:

- Setting temperature=0.0 ensures **precise, factual responses** (e.g., legal documents).
- Setting temperature=0.9 enables **creative storytelling** (e.g., marketing copy).

Parameter	Description	Effect on Output
temperature	Controls randomness (0.0 = deterministic, 1.0 = highly creative).	Lower values ensure predictable responses, higher values increase variability.
top_p	Nucleus sampling – sets diversity by filtering top probability tokens.	Lower values (e.g., 0.5) create focused answers , higher values (e.g., 0.9) allow more variety .
max_tokens	Limits response length.	Prevents excessively long outputs, useful for API efficiency.
frequency_penalty	Reduces repetition of words/phrases.	Higher values discourage repeated content in responses.
presence_penalty	Encourages introducing new topics.	Higher values make the response more exploratory and broad.



Key Takeaway:

Different parameters adjust precision, randomness, verbosity, and repetition.

Parameter Settings for Different Tasks

14


Use Case	Recommended Parameters
Fact-Based Q&A	temperature=0.0, top_p=0.5, max_tokens=150
Creative Writing	temperature=0.9, top_p=1.0, max_tokens=500
Summarization	temperature=0.3, top_p=0.8, max_tokens=250
Conversational AI	temperature=0.7, top_p=0.9, max_tokens=150, presence_penalty=0.6
Structured Data Extraction	temperature=0.1, top_p=0.5, max_tokens=100

Key Takeaway:

Tuning parameters aligns the model output with specific needs (accuracy, creativity, conciseness, etc.).

Effects of Changing Temperature & Top-P

15

 **Scenario:** Asking OpenAI: "Tell me an interesting fact about space."

Parameter Setting	Expected Response
temperature=0.0	"The speed of light in a vacuum is approximately 299,792 kilometers per second." (Strict factual answer)
temperature=0.9	"Did you know that a day on Venus is longer than a year on Venus?" (More unusual/interesting fact)
top_p=0.5	"There are over 100 billion galaxies in the observable universe." (More direct fact)
top_p=1.0	"Some scientists believe there might be unknown forms of life thriving in the extreme conditions of exoplanets!" (Diverse and speculative fact)

 **Key Takeaway:**

Higher temperature and top_p values increase creativity, while lower values improve precision.

Fine-Tuning LLM Outputs Effectively

16

✓ **Start with Defaults & Adjust Gradually**

- OpenAI defaults work well for most cases; tweak based on observed behavior.

✓ **Use Low temperature for Factual Tasks**

- Set temperature=0.0 for reliability (e.g., financial reports, legal summaries).

✓ **Limit Response Length (max_tokens)**

- Helps manage API costs and avoids excessively verbose responses.

✓ **Use frequency_penalty to Reduce Repetition**

- Good for avoiding redundant answers in long outputs.

✓ **Balance temperature and top_p**

- Avoid setting **both high**—use one or the other to fine-tune randomness.

Confidential & Proprietary – Not for Distribution

Demo



Lab



