



Integrating Vector, Relational, and JSON Data

1. Overview of AI Workflows	2. Why Postgres as a vector store	3. Storing and managing vectors	4. Querying the vector store
<ul style="list-style-type: none"> Look at high-level architecture - LLMs, vector stores and JSON Look at key vocabulary and concepts (embeddings, vectors, hybrid queries, etc.) 	<ul style="list-style-type: none"> What is a vector store? Key concepts and use cases. Why Postgres and how does it compare with other market tools Setting up Postgres with vector capabilities (pgvector) Lab: Install and configure Postgres using Docker 	<ul style="list-style-type: none"> Generating embeddings: Overview of tools and workflows Storing and organizing embeddings in Postgres Strategies for handling large datasets including chunking Dense and sparse vectors Lab: Generate embeddings for a dataset and store them 	<ul style="list-style-type: none"> Techniques for similarity search: k-NN, cosine similarity Using indexes to optimize vector queries Reranking results Lab: Query stored vectors to retrieve similar items (document/image search)
5. Querying LLMs with retrieved data	6. NoSQL with JSON in Postgres	7. Integrating Vector, Relational and JSON Data	8. Putting it all together
<ul style="list-style-type: none"> Recap on querying LLMs vis APIs Best practices for combining vector retrieval with LLM prompts Prompt configuration parameters (temperature, top-k, etc) Lab: Build a pipeline where vector store results enhance LLM responses (context-aware Q&A, etc) 	<ul style="list-style-type: none"> Overview of JSON/JSONB support in Postgres Querying JSONB data with SQL Indexing JSONB data for performance Lab: Design a schema mixing vector, relational and JSONB data for a sample project 	<ul style="list-style-type: none"> Building hybrid queries to power advanced workflows Case study: Combining embeddings, metadata (relational) and configurations (JSON) Lab: Implement a hybrid query to support a sample AI use case 	<ul style="list-style-type: none"> Full stack pipeline demo: Retrieve data, query the LLM and return results Debugging and optimising the workflow Spotlight on LLM frameworks Lab: Build a working application combining all elements



Plan for the session

- Understand the Benefits of Hybrid Queries
- Write SQL Queries that Combine Vector, Relational, and JSONB Data

Leveraging Multiple Data Types for AI-Powered Workflows

4

📌 Key Points:

- AI applications require **multiple types of data**:
 - **Relational** (structured metadata like categories, timestamps, user data).
 - **Vector embeddings** (semantic meaning and similarity).
 - **JSONB** (flexible, semi-structured configurations and additional attributes).
- Hybrid queries allow **rich, context-aware data retrieval**.

📌 Key Takeaway:

Hybrid queries improve AI workflows by combining structured, unstructured, and semantic search capabilities.

Breaking Down the Query Stack

5

◆ 1 Relational Data (PostgreSQL Table Columns)

- Stores **structured fields** (e.g., id, name, created_at).
- Enables **fast lookups and filtering** using indexes.

◆ 2 Vector Data (pgvector for Embeddings)

- Stores **semantic representations** of text/images.
- Enables **nearest neighbor search** for similarity matching.

◆ 3 JSONB Data (Flexible Attributes & Configurations)

- Stores **semi-structured data** (e.g., {"price": 29.99, "stock": 100}).
- Allows **flexible querying and updates** without schema changes.

📌 Key Takeaway:

Each data type serves a different purpose, and combining them makes AI retrieval more powerful.

Query Flow: Retrieving Context-Rich Data

6

1 User Query → Generate Vector Embedding

- Convert text query into an embedding using bge-m3.

2 Vector Similarity Search → Retrieve Relevant Entries

- Use pgvector to find the most similar stored items.

3 Relational Filtering → Narrow Down Results

- Apply filters (e.g., WHERE category = 'programming').

4 JSONB Extraction → Enrich Data with Additional Fields

- Fetch relevant metadata (e.g., price, configuration, availability).

Demo



Real-World Scenarios That Benefit from Hybrid Data

8

✓ Semantic Search + Metadata Filtering

- Find **similar research papers** but filter only by **peer-reviewed publications**.

✓ Personalized Recommendations

- Retrieve **movies similar to "Inception"**, then filter by **user preferences (JSONB settings)**.

✓ AI-Powered Knowledge Retrieval

- Find **similar troubleshooting issues**, but **prioritize ones from recent cases**.

Optimizing Multi-Modal Queries in PostgreSQL

9

✓ Use Indexing for Performance

- **GIN indexes** for JSONB
- **IVFFLAT or HNSW indexes** for pgvector

✓ Limit Query Scope for Efficiency

- Use LIMIT, WHERE, and ORDER BY to **reduce search space**.

✓ Normalize Data When Necessary

- Store **high-frequency relational fields** separately to **avoid JSONB bloating**.

✓ Cache Query Results for AI Pipelines

- Precompute and store results for **frequent queries** to improve response times.

Lab



Apply Hybrid Querying in a Real-World AI Use Case (*Hands-on Lab*)

- Design a **real-world query** that fetches relevant **vector-based** results, filters them using **relational metadata**, and **enriches responses with JSONB fields**.
- Implement a **case study** demonstrating **multi-source AI retrieval** (e.g., **retrieving books, filtering by metadata, and adjusting configurations dynamically**).

